

Arrays (JFS 13)

- Motivation - bulk data handling and processing
- Declaring and using
- Basic algorithms
- More advanced techniques

Arrays: Motivation (JFS 13)

- Suppose a meteorologist needs to total rainfall over a period for each day of the week
- By convention, say:
0 = Monday, 1 = Tuesday, ..., 5 = Saturday, 6 = Sunday
- The rainfall for each day is to be recorded as a floating-point number in mm
- The total rainfall per day could be stored in individual variables:

```
float rain0, rain1, rain2, rain3, rain4,  
      rain5, rain6;
```

- This is clumsy, and would become unworkable if the rainfall were to be recorded for *every day of a year!*
- It needs a collection of identical kinds of items
 - Each identified by a particular index or subscript
 - For example the day number, size code
- Collections of identical items occur frequently, for example:
 - sales of products, indexed by code 0, 1, 2, ...
 - diameters of bolts, indexed by size 0, 1, 2, ...
 - marks for a test, indexed by mark 0, 1, 2, ...

- For such a situation, Java provides the **array**
 - A collection of identical kinds of elements
 - Indexed by a subscript
- An array with a single subscript is called a list, vector or one-dimensional array
- The number of items in an array is fixed by its declaration
 - The array cannot grow or shrink
- [The libraries contain other data structures that *can* grow and shrink.]

Declaring Arrays

- An array is created in two steps:
 - 1) A variable is declared to hold the array:
`float rain[]; // name an array of floats`
 - 2) Space is allocated for the array:
`rain = new float[7];`
- Both steps are needed
 - Until space has been allocated an array is only just a name
- It is quite common to do both steps at the same time:
 - `float rain[] = new float[7];`
- [If a global variable then also (usually) `private`.]
- Note: Since an array is *non-primitive data*, the array variable `rain` actually holds a reference/pointer to the allocated memory

- Notice the use of brackets
 - They indicate that a subscript will be used
- Actually the brackets can be written after the type or name when an array is declared:


```
float rain[];
float[] rain;
```
- The first form is closer to how subscripts are used, though the second form is used in [Bell and Parr]
- An array like `rain` can hold 7 values
- The size of an array is fixed once space has been allocated
 - Can refer to it in Java code like this: `rain.length`

rain →	1.5	3.2	0.0	4.0	7.1	6.9	5.3
Day Number →	0	1	2	3	4	5	6

**Subscript
or Index**

- As shown in the figure, array subscripts start at 0
 - The first element is numbered 0
 - The last is numbered 6
- It is easy to mistakenly expect such an array to be numbered from 1 to 7 since it has 7 values
- A run-time error will occur if the program tries to refer to a subscript is out of range
 - Here, less than 0 or greater than 6
- Once an array has been declared and space has been allocated, the values of its elements can be set and used

Using Arrays

- Array elements can be used directly:


```
rain[5]           // Rainfall for Saturday (5)
rain[0] + rain[1] + rain [2]
// Total rainfall for Mon - Weds
```
- "Indexed variables" can be used in expressions and on the left of assignment statements:


```
rain[4] = 6.2; // Set rainfall for Friday
float mtwTotal = rain[0] + rain[1] + rain[2];
- Just like any variable
```
- The subscript or index can be an *expression*:


```
int d = 3;      // d is a "day of interest"
rain[d] = 4.7;
total = rain[d-1] + rain[d] + rain[d+1];
- A subscript expression is re-evaluated at every access
```

- Since arrays contain identical kinds of elements, they are particularly suitable for *repetitive* operations

- Often a **for** loop is appropriate for this

- e.g. to add up all the elements in an array:

```
float total = 0.0f; // For week total
for (int day = 0; day < rain.length; day++) {
    // go through days
    total = total + rain[day]; // add one day
}
```

Note: Refers to a different element of rain at each repetition of loop body

- To calculate the average rainfall, calculate the total and then do:

```
float average = total / rain.length;
```

- Note the use of **rain.length**

- The code does *not need adapting* if we change the array size

Declaring Arrays (again)

- Since the programmer must decide the size of an array in advance, it is good practice to use a constant for this
- A constant is a 'variable' declared as **final**; the value given to it cannot be changed
- As an example of a constant, the size of an array can be declared and used as:

```
final int days = 7; // The number in a week
float rain[] = new float[days];
// Use this as the size
```

- This has several advantages:

- The meaning of the array size might not be obvious:
e.g. there might be different arrays of size 7
- giving the size a *name* makes the *intention* clearer
- There might be several arrays that *must have the same size*:
redefining one constant will affect *all sizes consistently*

Initialising Arrays

- The elements of a new array will probably contain nothing useful
 - So the program must usually give each item a specific value initially
- If the array contains numbers, a good initial value is probably zero:

```
rain[0] = 0.0f; rain[1] = 0.0f;
rain[2] = 0.0f; rain[3] = 0.0f;
rain[4] = 0.0f; rain[5] = 0.0f;
rain[6] = 0.0f;
```

- However this needs one assignment per array element
 - Not a good idea if the array has many elements!
- A much better idea is to use a *loop* to carry out such a repetitive task:


```
for (int day = 0; day < days; day++) // Week
    rain[day] = 0.0f; // Initialise one day
```
- Notice that the number of **days** has been used here as a constant, instead of writing 7 into the program
- This makes it clearer that the loop deals with all array elements
- Using 7 literally would be problematic if the size of array were changed later (perhaps not in this application!)

- A loop is also useful if the initial values can be decided according to some rule
- Consider an array that stores the 12-times table for multiplier 0 to 10:
- The initialisation might be:

```
int twelveX[] = new int[11];
for (multiplier = 0; multiplier < 11;
     multiplier++) {
    twelveX[multiplier] = multiplier * 12;
}
```

twelveX →	0	12	24	36	...	108	120
Subscript →	0	1	2	3	...	9	10

- Sometimes the initial values are irregular and have to be given directly:


```
// ages of children 0, 1, 2
age[0] = 2; age[1] = 3; age[2] = last + 4
```
- If these values are to fill the whole array, can conveniently use an *array initializer*:


```
int age[] = {2, 3, last + 4};
```
- Note that the size of such an array is implied by the number of initial values
 - And it replaces the declaration, space allocation and initial assignments
- Other ways of getting initial values into an array:
 - Ask the user to type them into a text field in a window, then read and store these values
 - Read the values from a file on disc

Using Arrays – more algorithms

- It is also often necessary to search through an array according to some criterion, e.g. find the day with the highest rainfall:

```
int wettest = 0;      // Assume day 0 the wettest
for (int day = 0; day < rain.length; day++) {
    // Go through the days
    if (rain[day] > rain[wettest]) // Current wetter?
        wettest = day;           // Yes, note it
}
```

- At the end of this loop, a value for `wettest` like 4 means that day 4 had the highest rainfall (or equal wettest)
- This is a *common pattern*:
 - Wettest, driest, number of very wet/quite dry days, find a completely dry day, and many other situations

- How many quite dry days?

```
int dryDays = 0;      // To count the days
for (int day = 0; day < rain.length; day++) {
    // Go through the days
    if (rain[day] < 1.0f) // Quite dry day?
        dryDays++;         // Yes, count it
}
```

- Sometimes it is necessary to search through an array for a *particular value*, e.g. is there a day that had no rain?
 - Special case: There might be no such day

```
int dryDay = -1;           // Day initially unknown
for (int day = 0; day < rain.length; day++) {
    // Go through the days
    if (rain[day] == 0.0f) // No rain this day?
        dryDay = day;     // Note this as a dry day
}
```

- So if there were no rain on day 5, `dryDay` would be set to 5
 - And if there is no dry day, then `dryDay` remains -1

- There are two potential complications here:
 - If there was no rain on days 3 and 5, `dryDay` would be set to 5 (i.e. the *last* such case)
 - If *every* day had some rain, `dryDay` would stay at -1; a later part of the program would have to check for this condition
- When searching an array, often the *first* suitable value is wanted
 - Continuing to search the array wastes time (especially if it is large)
- A `for` loop can be undesirable because it searches *all* the array elements

Short-cutting searches: two solutions

- First solution:
- Instead a **while** loop may be better because it can finish as soon as the required value is found:

```

int dryDay = -1;           // Day initially unknown
int day = 0;               // Start with first day
while (dryDay == -1 &&      // Still looking and ...
       day < rain.length) { // not reached array end?
    if (rain[day] == 0.0f) // No rain this day?
        dryDay = day;     // Note this as a dry day
    day++;                // Move to next day
}

```

- Notice that there are *two* reasons that such a loop terminates:
 - Finding the desired condition
 - or reaching the end of the array
- Also note that a loop like this has to explicitly advance the loop counter (**day**)
 - Forgetting to do this will mean the program loops indefinitely!

- Second solution:
- We can use a **for** loop with a **break;** statement:

```
int dryDay = -1;           // Day initially unknown
for (int day = 0; day < rain.length; day++) {
    // Go through the days
    if (rain[day] == 0.0f) { // No rain this day?
        dryDay = day;      // Note this as a dry day
        break;
    }
}
```

- Executing **break;** causes an immediate exit from the *closest enclosing loop*:
 - Best of both worlds?
 - Some programmers disagree with using **break;**

Arrays as Parameters

- Just as a method can take simple types like integers or floats as parameters, it can take any kind of object including arrays
- This is where **array.length** is particularly useful, because the method cannot know in advance how long the array will be
- For example, a general method can be written to set an array of floats to zero:

```
// Take any float array, set elements to 0
private void zeroise(float arr[]) {
    for (int s = 0; s < arr.length; s++)
        // Go through all elements
        arr[s] = 0.0f; // Set element s to zero
}
```

- Another method might copy the contents of one array to a second array if they are the same length:

```
// Copy array1 to array2 if compatible
private void copy(float array1[], float array2[]) {
    int len1 = array1.length; // First array length
    int len2 = array2.length; // Second array length
    if (len1 == len2) {      // Arrays same length?
        for (int s = 0; s < len1; s++)
            // Go through all elements
            array2[s] = array1[s]; // Copy element s
    }
}
```

- Notice that this is *different* from assigning the first array to the second:

```
anotherArray = anArray;
- which makes anotherArray the same array as anArray:
- anArray and anotherArray become pointers to the same
  memory elements
```

- In contrast:

```
copy(anArray, anotherArray);
- gives two entirely independent copies of the same values
- Copying keeps the arrays separate
```

Other Kinds of Arrays (JFS 13, 14)

- Arrays are not restricted to holding numbers:

```
char vowels[] = {'a', 'e', 'i', 'o', 'u'};
String name[] = {"Joe", "Sue", "Pat"};
boolean married[] = {false, true, false};
```

- Perhaps `name` and `married` are "parallel arrays", holding corresponding information
- In fact arrays can hold *any kind of object*, such as:
 - Records - see example on next slide
 - Graphical items like buttons:

```
JButton[] controls = new JButton[10];
```

- Given the definition of a telephone directory entry earlier, a whole telephone directory might be declared and set up as:

```
Entry directory[] = new Entry[1000];
directory[0] =          // Set first subscriber
    new Entry("Jo Smith", "467420", true);
directory[1] =          // set second subscriber
    new Entry("Mary Brown", "467423", false);
...

```

- Obviously, to fill in the complete directory would be a lot of work
 - Normally such information would be read from a file
- Working with the entry data can look complex:

```
directory[i].setUnlisted(true);
g.drawString(directory[i].getName(), x, y);
```

- Listing the whole (listable) directory in text area `output`:

```
output.setText("");
for (int n = 0; n < directory.length; n++) {
    if (!directory[n].isUnlisted()) {
        output.append(directory[n].getName() +
": " + directory[n].getNumber() + "\n");
    }
}
```

- Searching for the number for a subscriber:

```
String name = nameText.getText(); // To be found
int number = -1; // Number initially unknown
for (int n = 0; n < directory.length; n++) {
    if (name.equals(directory[n].getName())) {
        number = directory[n].getNumber();
        break;
    }
}
```

- Note:

- `.equals` requires a precise match
- `.equalsIgnoreCase` is more flexible
- Partial match would be even better

Two-dimensional Arrays

- The elements of an array can also be arrays!
- This is best thought of as a table, also called a two-dimensional array
- Examples of two-dimensional arrays are:
 - A contour map, where each pair of x and y coordinates gives a height
 - A chess-board, where each square gives the name of a piece occupying it
 - Sales in a chain of retail stores, where one subscript identifies the shop and a second subscript identifies the product code; array elements give the sales of each product per store

- The contour map might be declared as:


```
final int EASTINGS = 1000;    // Grid dimensions
final int NORTHINGS = 1000;
float height[][] =
    new float[EASTINGS] [NORTHINGS];
```

 - **EASTINGS** is the west-east coordinate
 - **NORTHINGS** is the south-north coordinate
 - **height** is in metres - one per x, y coordinate
- Heights in the contour map would be set and used as follows:


```
height[5][3] = 21.5f;    // Set height at grid (5,3)
hilltop = height[4][6]; // Get height at grid (4,6)
```
- As can be seen, two-dimensional arrays are like one-dimensional arrays except that they have a pair of subscripts
- See next slide

- We can picture the two-dimensional array like this:

9												
8												
7												
6					97.1							
5												
4												
3						21.5						
2												
1												
0												
↑	North	East →	0	1	2	3	4	5	6	7	8	9

- Note: There would be a value in every element of the grid

- Nested **for** loops are often used to process two-dimensional arrays
- For example: If the retail chain sales table is declared like this:

```
final int STORES = 100;      // Size of chain
final int PRODUCTS = 1000; // How many stocked
int sales[][] = new int[STORES][PRODUCTS];
```

- (And the sales have all been recorded in the array)
- Then the total sales can be calculated by:

```
int totalSales = 0;
for (int s = 0; s < STORES; s++)
    for (int p = 0; p < PRODUCTS; p++)
        totalSales = totalSales + sales[s][p];
```

Array Example

Rainfall Array

End of section